

- Security Architecture
  - Security Architecture Overview
  - Authentication System
    - Development Environment
    - Production Environment
    - Authentication Flow
  - Authorization Model
    - JWT-Based API Authorization
    - Fine-Grained Access Control
    - User Permissions Model
  - Data Protection
    - Data-at-Rest Encryption
    - Data-in-Transit Encryption
    - Sensitive Data Handling
  - Network Security
    - API Protection
    - VPC Integration (Optional)
  - Secure Development Practices
    - Application Security
    - Dependency Management
    - Secure Configuration
  - Monitoring and Incident Response
    - Security Monitoring
    - Incident Response Capabilities
  - Compliance Considerations
  - Authentication Implementation Details
    - Frontend Authentication Service
    - API Security Integration
  - Security Best Practices for Developers

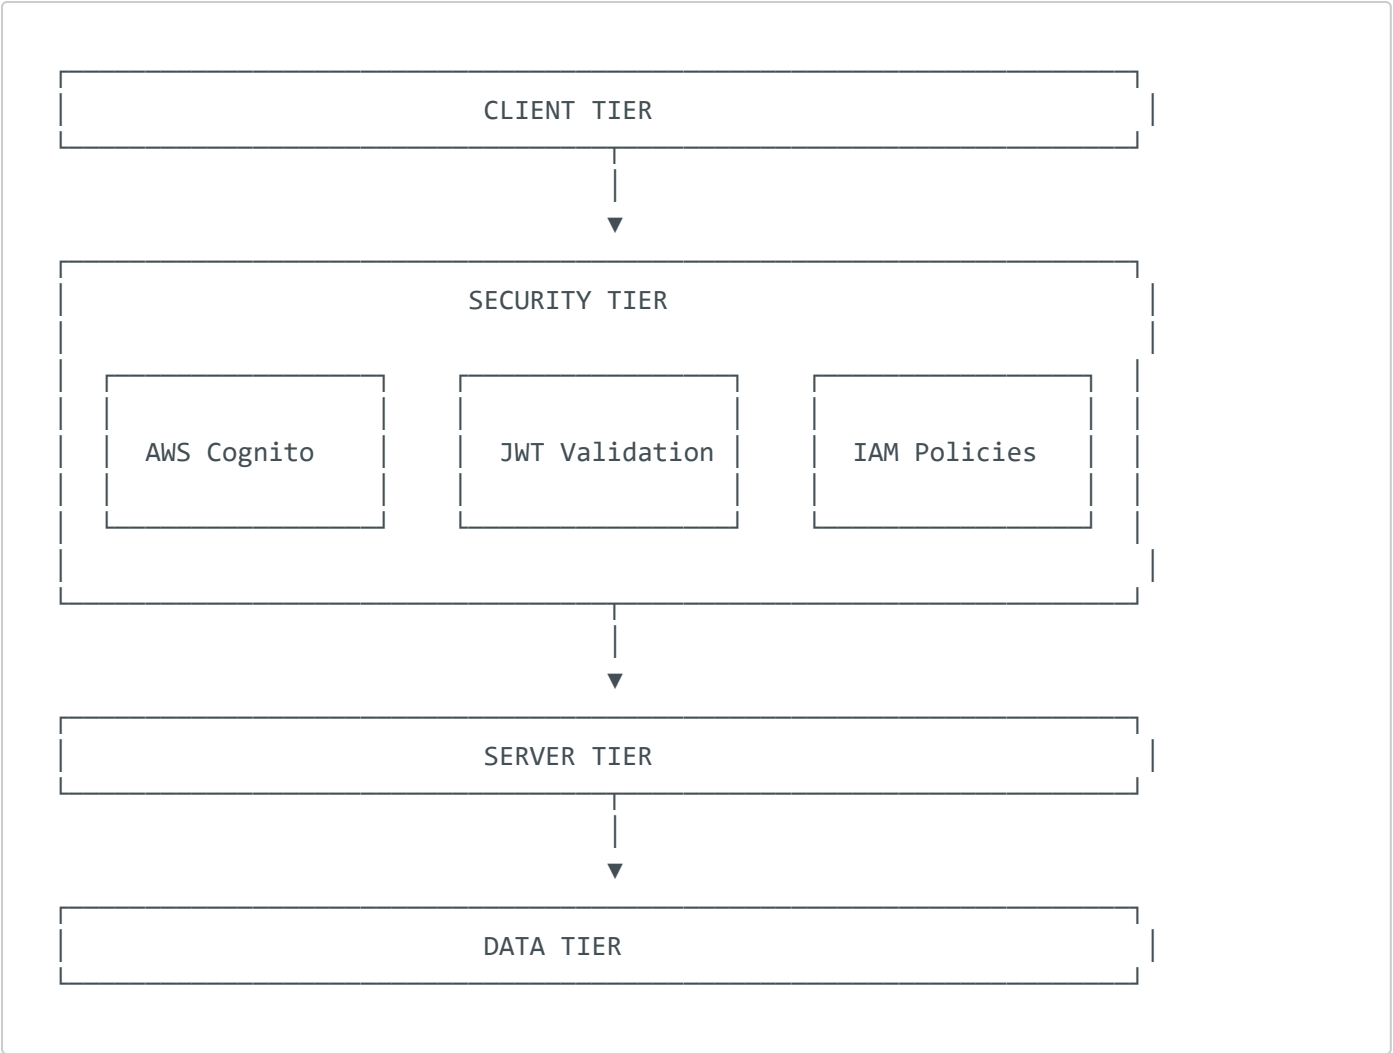
# Security Architecture

---

This document outlines the security architecture of the Document Processing Accelerator, detailing the approach to authentication, authorization, data protection, and secure communications.

# Security Architecture Overview

---



## Authentication System

---

The Document Processing Accelerator implements a dual authentication system:

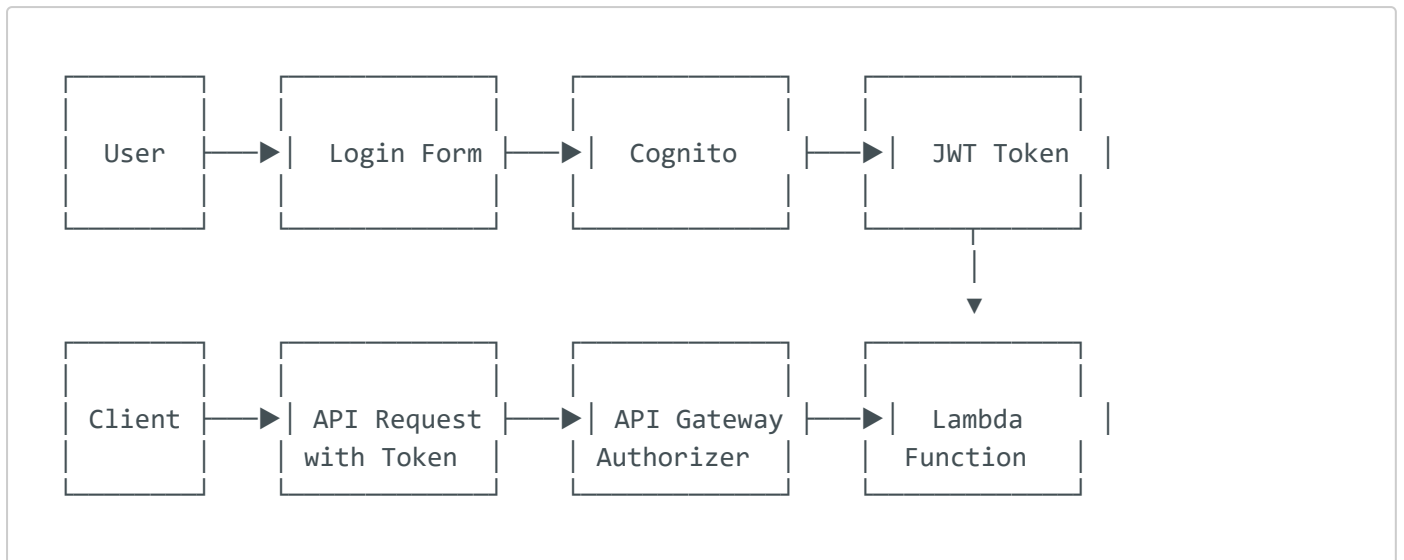
### Development Environment

- **Mock Authentication Service** (`mockAuthService.ts`)
  - Simulates authentication flows without external dependencies
  - Stores authentication state in local storage
  - Implements the same interface as the production authentication service

### Production Environment

- **AWS Cognito Authentication** (`amplifyAuthService.ts`)
  - Uses AWS Amplify v5 API for integration with Cognito
  - Supports user sign-up, sign-in, and password recovery
  - Implements MFA capability for enhanced security
  - Issues JWT tokens for API authorization

## Authentication Flow



## Authorization Model

### JWT-Based API Authorization

- Tokens issued by Cognito are validated at the API Gateway level
- Custom authorizer Lambda function validates token signatures
- Token claims determine user permissions

### Fine-Grained Access Control

- IAM roles and policies restrict Lambda function permissions
- DynamoDB conditional expressions enforce user-based data isolation
- S3 bucket policies limit access to user-specific paths

## User Permissions Model

Role	Document Access	Administrative Functions
Regular User	Own documents only	None
Admin	All documents	User management
Super Admin	All documents, system config	System configuration

## Data Protection

---

### Data-at-Rest Encryption

- **S3 Buckets:** Server-side encryption (SSE-S3) for document storage
- **DynamoDB:** Encryption enabled for all tables
- **Lambda Environment Variables:** Encrypted with KMS
- **Secrets:** Stored in AWS Secrets Manager with automatic rotation

### Data-in-Transit Encryption

- **API Gateway:** HTTPS enforcement for all endpoints
- **S3 Pre-signed URLs:** HTTPS with temporary credentials
- **Internal AWS Service Communications:** TLS 1.2+

### Sensitive Data Handling

- PII data is encrypted at the application level before storage
- Document content is processed in memory and not stored in plaintext
- Encryption/decryption operations use AWS KMS for key management

## Network Security

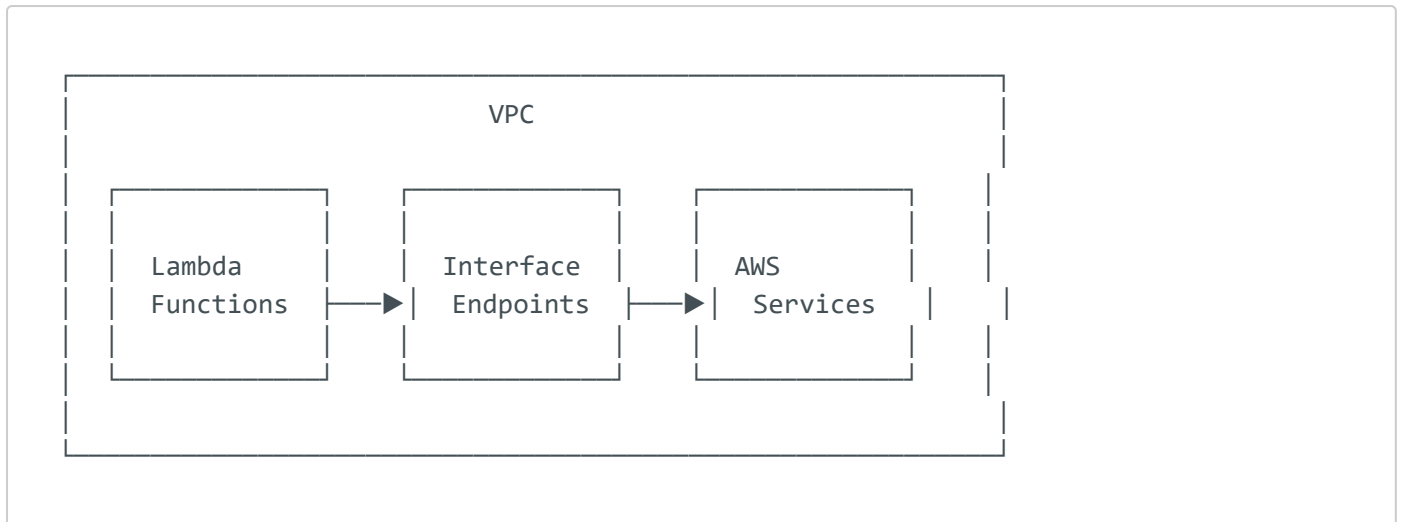
---

### API Protection

- API Gateway provides front-end protection from attacks
- Rate limiting prevents abuse and DoS attacks

- WAF integration blocks common attack patterns

## VPC Integration (Optional)



- Lambda functions can be configured to run within a VPC
- Private network paths to AWS services via VPC endpoints
- Security groups restrict traffic flow

## Secure Development Practices

---

### Application Security

- Input validation on both client and server
- Output encoding to prevent XSS attacks
- Proper error handling that doesn't leak sensitive information
- Content Security Policy implementation

### Dependency Management

- Regular dependency scanning for vulnerabilities
- Automated updates for critical security patches
- Dependency pinning to prevent supply chain attacks

### Secure Configuration

- Environment-specific configuration with proper separation
- No hardcoded secrets or credentials
- Infrastructure as Code with version control

# Monitoring and Incident Response

---

## Security Monitoring

- CloudWatch Logs for Lambda function monitoring
- CloudTrail for AWS API activity logging
- Custom alerting for suspicious activities

## Incident Response Capabilities

- Automated alerts for security anomalies
- Predefined response playbooks for common scenarios
- Regular security incident drills

## Compliance Considerations

---

The Document Processing Accelerator is designed with the following compliance frameworks in mind:

- GDPR: Data privacy and protection measures
- SOC 2: Controls for security, availability, and confidentiality
- HIPAA: If configured for healthcare document processing

## Authentication Implementation Details

---

### Frontend Authentication Service

The authentication service is designed with a provider pattern that allows seamless switching between development and production implementations:

```
// authServiceProvider.ts
import { AuthService } from '../types/auth';
import { mockAuthService } from './mockAuthService';
import { amplifyAuthService } from './amplifyAuthService';

// Determine which auth implementation to use
const useRealAuth = process.env.REACT_APP_USE_REAL_AUTH === 'true';

// Export the appropriate auth service implementation
export const authService: AuthService = useRealAuth
  ? amplifyAuthService
  : mockAuthService;
```

## API Security Integration

The API client automatically includes authentication tokens in requests:

```
// apiClient.ts
import { authService } from './authServiceProvider';

export const apiClient = {
  async get(endpoint: string) {
    const token = await authService.getToken();

    return fetch(`${API_BASE_URL}${endpoint}`, {
      headers: {
        'Authorization': `Bearer ${token}`,
        'Content-Type': 'application/json'
      }
    }).then(handleResponse);
  },

  // Additional methods (post, put, delete) follow the same pattern
};
```

## Security Best Practices for Developers

When working on the Document Processing Accelerator, developers should:

1. **Never disable authentication** in production environments
2. **Always validate input** from users and external systems
3. **Use parameterized queries** when accessing DynamoDB
4. **Follow the principle of least privilege** for IAM roles

5. **Keep dependencies updated** to mitigate security vulnerabilities
6. **Enable MFA** for their AWS console access
7. **Use pre-signed URLs** for all S3 operations exposed to end users